


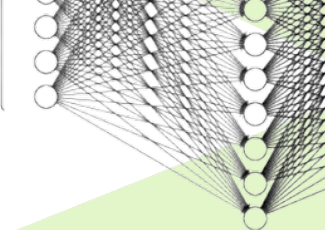
Neural Networks: Classical vs Quantum


Lourens van
Niekerk





MNIST dataset

- (Artificial) neural networks simulate the learning mechanism of biological neural networks.
 - Components:
 - Neurons (units) with weights that scale their input
 - Synapses (connections) with weights that distribute neuron output as input to other neurons
 - Mechanism:
 - NN propagates input data through it, into output
 - Weights in NN adjust over iterations to via incentives
 - Incentives can be rewards, example matching or patterns
 - → Reinforcement, Supervised or Unsupervised Learning
- 
- A diagram of a neural network structure. It shows multiple layers of nodes (represented by circles) connected by lines (synapses). The nodes are arranged in a grid-like pattern, with lines connecting them in a complex, overlapping manner. The diagram is positioned in the top right corner of the slide.



Key terminology



Features

Epochs

Noise

Batch size

Learning rate

Backpropagation

Activation function

Regularization

Loss function

Classification & Regression

Barren plateaus




Features, Epochs & Noise

- Features are the inputs to the neural network
 - Could be raw data or transformations
- Epoch = iterations
 - How many full processes on the (NN) model occurred
 - I graduated high school after 18 epochs with a batch size of ~365 days (← my cost increased suddenly)
- Noise is how much misinformation / misinterpretation / external factors are affecting the data
 - Here it only refers to noise in the data
 - This is not noise affecting quantum operations

500 data points

Blue and Orange
have spies





Epoch
000,000

Learning rate
0.001

Activation
ReLU

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 70%

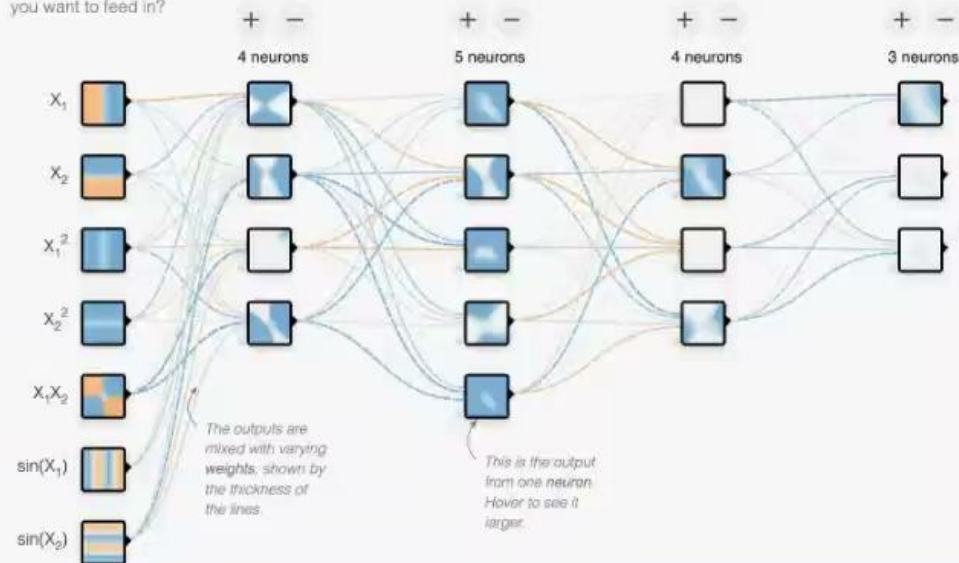
Noise: 0

Batch size: 15

REGENERATE


FEATURES

Which properties do you want to feed in?



OUTPUT

Test loss 0.520
Training loss 0.522



Colors shows data, neuron and weight values.

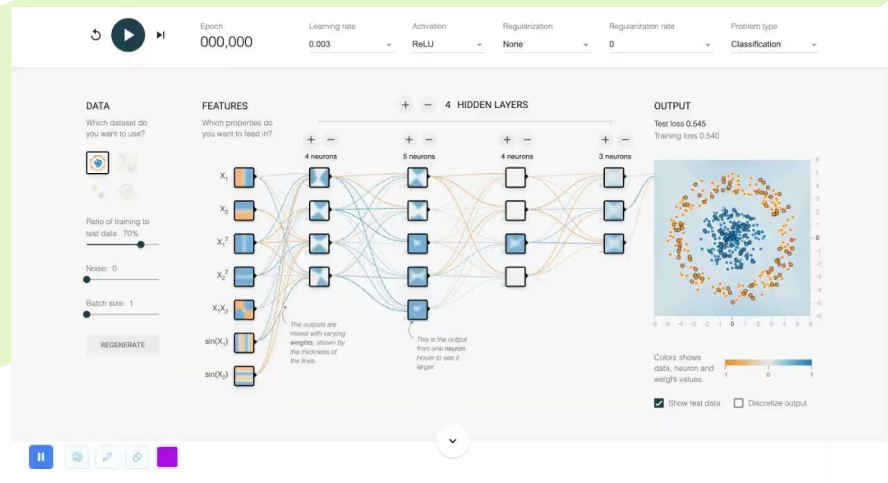



☒ Show test data ☐ Discretize output

Batch size

- Determines how many samples used per forward & backward pass
- After each batch, model makes predictions, and adapts weights
- Different types
 - Batch Gradient Descent (BGD)
 - All training samples
 - Mini-Batch Gradient Descent (MBGD)
 - Somewhere in-between
 - Stochastic Gradient Descent (SGD)
 - Sample of 1

Understanding Batch Size in Neural Network



Learning rate

- How much do we allow the model to adjust after each epoch

- Too low: forever to train
- Too high: unstable

- Useful to have adaptive learning rate → ADAM

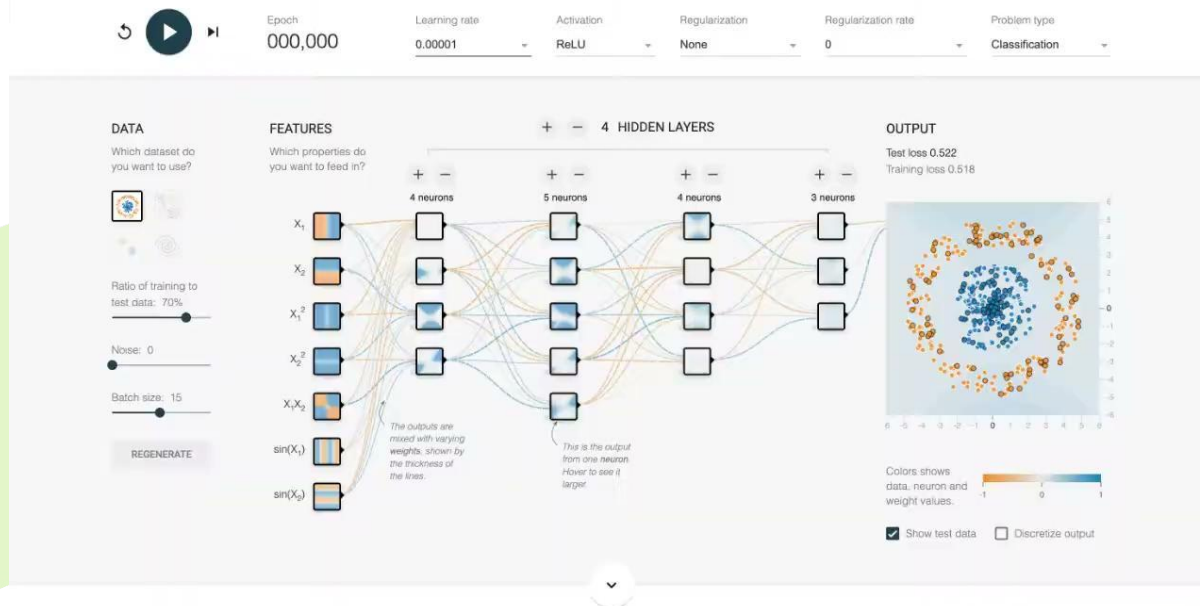
- [Schedules of decay](#)

$$\omega_{i+1} = \omega_i - \lambda \cdot \nabla L(\omega_i)$$

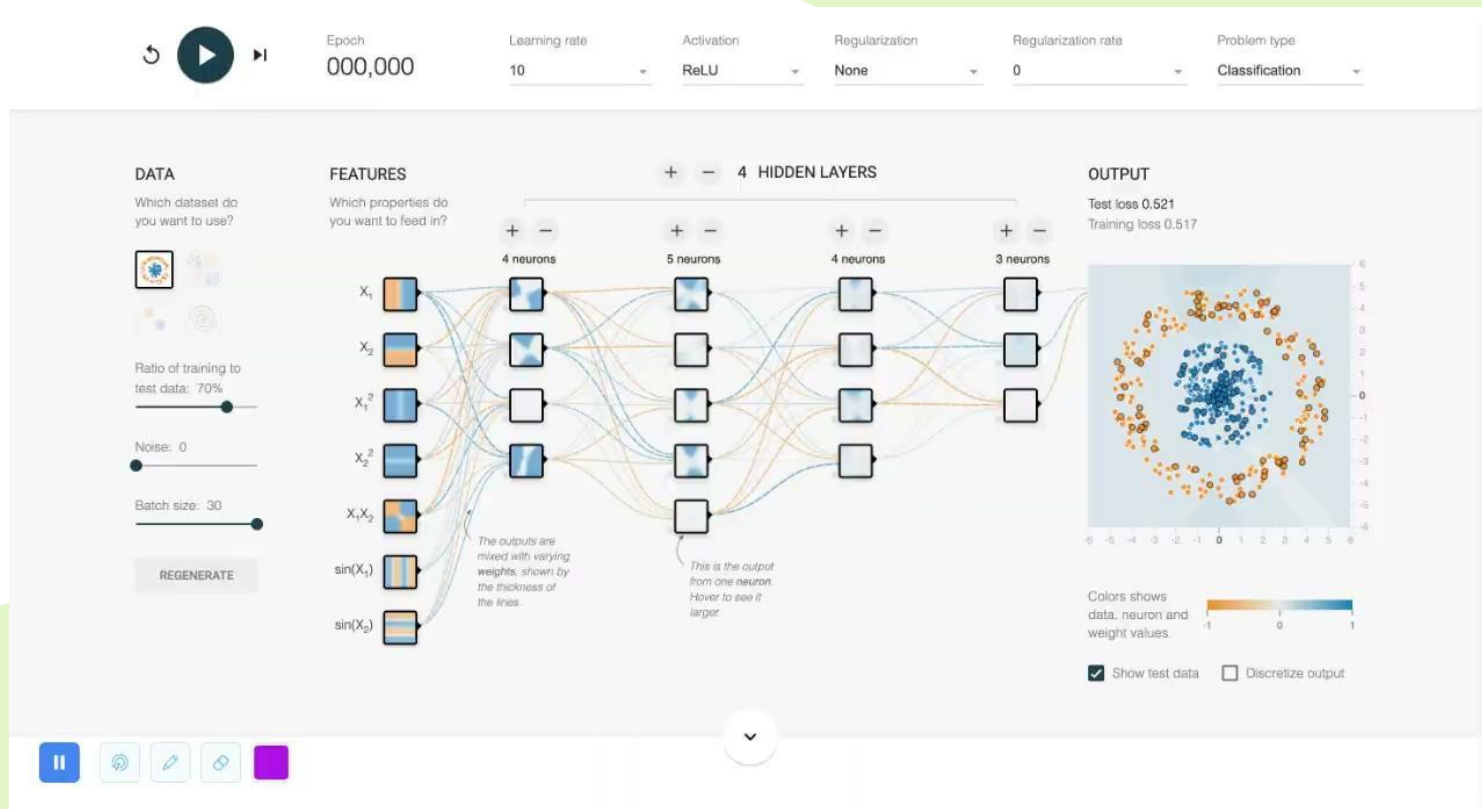
- ω_i = weights of model

- λ = learning rate

- $\nabla L(\omega_i)$ = gradient of loss function



Unstable learning rate example



Loss function

A medium to facilitate optimization of model outputs (also called cost function)


Often incentive is minimization

Examples:


- Mean Square Error (MSE) ← sensitive to outliers
- Mean Absolute Error (MAE)
- Binary cross-entropy (Log loss)
- Kullback-Leibler Divergence loss (KL Divergence)
 - Great for probability distributions

Activation function


- A function applied to neuron output
- Introduces nonlinearity to the model
- Allows classification of real-world complex data
- Perceptron just uses Sign
- ReLU: Rectified Linear Unit
- ReLU favorite, with many variants




(a) Identity




(b) Sign




(c) Sigmoid



(d) Tanh



(e) ReLU



(f) Hard Tanh

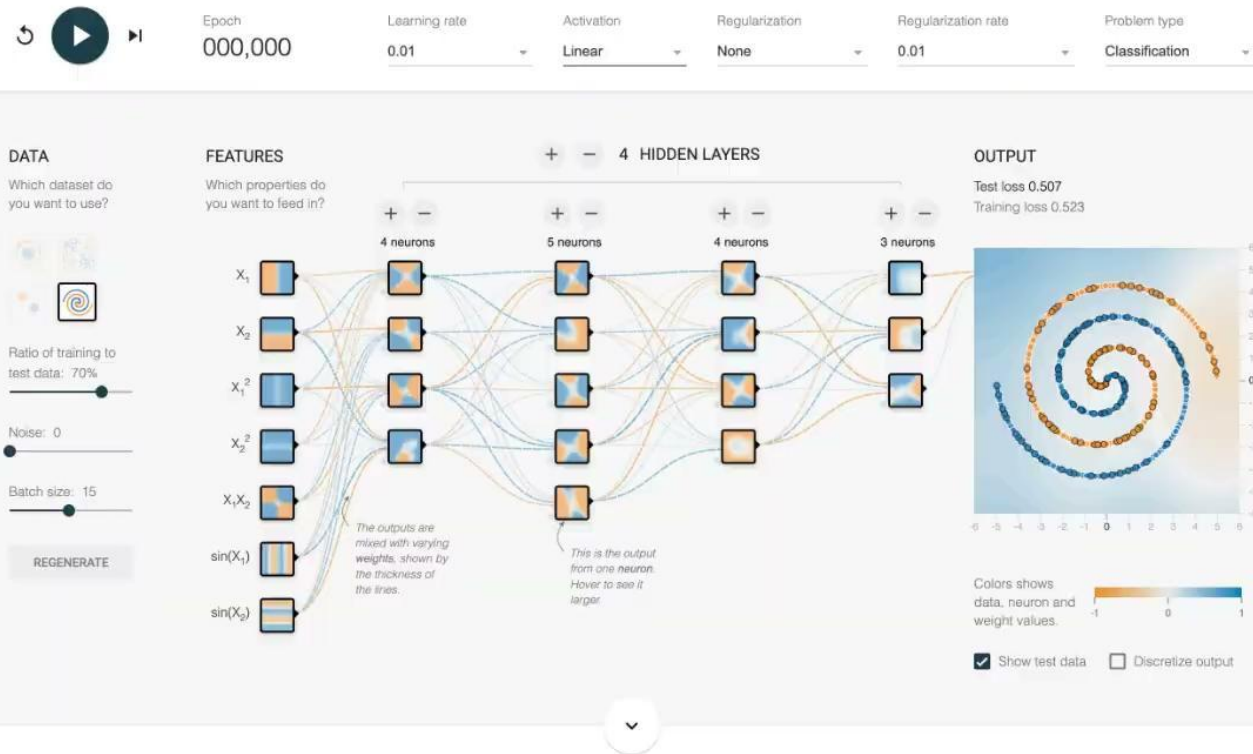
Credit: Aggarwal

$$ReLU(x) = \max(0, x)$$

$$Sigmoid(x) = \frac{1}{1+e^{-x}}$$

Vanishing gradient problem

Activation function



Data can be
obviously
impossible to
separate
linearly

Linear activation function

Data could be easy to separate sometimes...




VS



Model: “No
problem!”


The problem



Things are bit more complex though...

← Model

Regularization

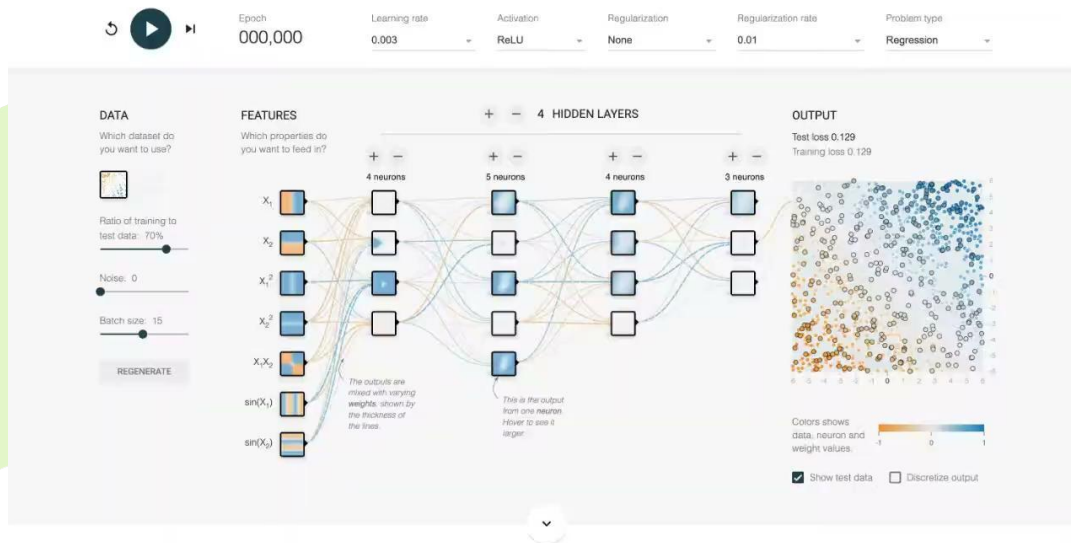


Underfitted

Good Fit/Robust

Overfitted

- Methods to counter for overfitting
 - Perfect fit on training data causes poor fit on test data
- Dropout: Drops the learning rate when static for an amount of epochs
 - Also good for barren plateaus
- L1 - Lasso regression
 - + absolute magnitude of feature coefficient
- L2 - Ridge regression
 - + square magnitude of feature coefficient



Backpropagation

ω_{ij} = Weight from neuron i to j

λ = Learning rate

e_i = Error at neuron i

$A_j(x)$ = Activation function at neuron j

$E = l(a_{out})$ = Output error / cost

Error after forward pass used to adjust weights in model backwards

Steps:

$$a_j = \sum_{\omega_{kj} \forall k} \cdot \omega_{ij} \cdot A_j(o_i)$$

- 1) Recursively assign neuron errors backwards to previous layers via weights

$$e_i = a_i (1 - a_i) \omega_{ij} \cdot e_j ; e_{out} = a_{out} (1 - a_{out}) E$$

- 2) Recursively calculate weight error gradient via neuron errors

$$\omega_{ij}^E = \frac{\partial E}{\partial \omega_{ij}} = e_j \cdot \lambda \cdot A_j(o_i)$$


- 3) Combine weights with weight error gradients

$$\omega^{new} = \omega + \omega_{ij}^E$$


Classification vs Regression

Discrete modelling vs Continuous modelling




“Which color is this?” vs “what value is this?”



VS




Barren plateaus



Credit: Kashif et al.

- A space where gradient descent-like approaches struggle
- Particularly a problem in quantum, as McClean et al. pointed out ([2018](#))
 - Variance in VQCs exponentially decrease to 0 as the qubits and/or layers increase
 - Make claim due to 2-design of random circuits
- Random quantum circuits are approx. 2-designs - [Harrow and Low \(2009\)](#)
 - Random gates over time yield a (Haar-distributed) unitary.
 - A 2-design has its variance equal to the Haar distribution.



Credit: McClean et al.

**“If you have a barren plateau,
all hope of quantum speedup
or quantum advantage is lost”**
- Marco Cerezo (Los Alamos, 2021)

Barren plateaus


Example:

**Find the (best)
watering hole**

- Roam around aimlessly and patiently
- Awareness without impulsiveness for descents
- Use the right tools:
 - Improved initialization → [random init. in PQCs](#)
 - Adaptive learning rates → [ADAM optimizer](#)
 - Regularization → [entanglement regularization](#)
 - Random walks → [quantum walk search](#)



A long neck helps...



Barren plateaus (BPs)

Reviews and surveys about BPs in VQCs: [2024](#), [2025](#)

Techniques to reduce BPs:

- Shallow circuits
- Iteratively change circuit ← Variable structured QNNs
- Alternative initialization strategies
 - Classical pre-training, structured random init., [parameter transfer](#), [initialization ansatz](#)

Techniques that are false promises:

- Changing optimizer ← can still require exponentially many shots
- Error mitigation in noise-induced BPs ← exponentially many resources

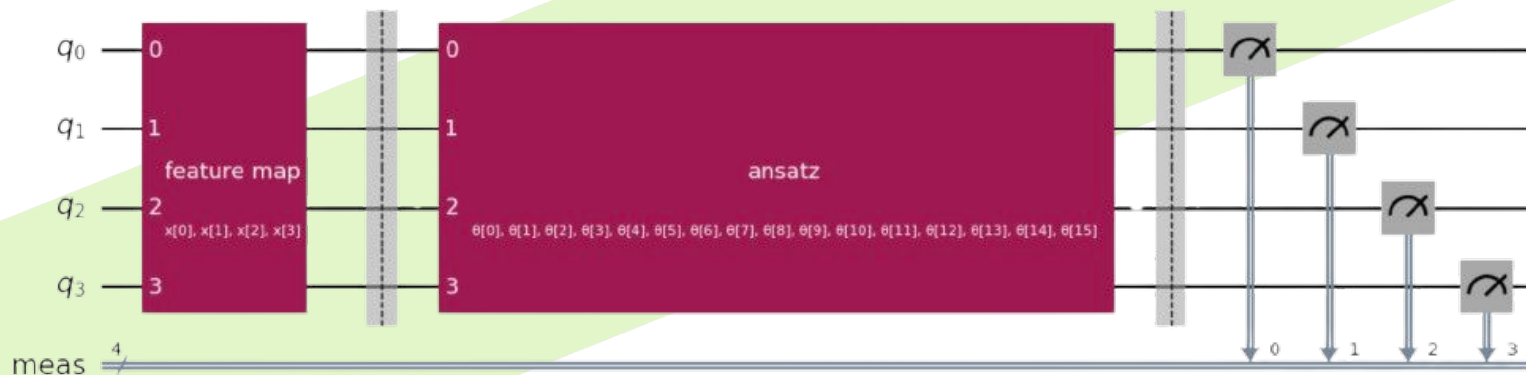
Quantum Neural Networks

Parameterized Quantum Circuits (PQCs)

Popular structure is VQAs, which is only what we will consider



Other structures include:

- [Continuous-Variable architecture \(CV\)](#)
- [Repeat-Until-Success \(RUS\)](#)




QNN changes to NN

- Entanglement layers ~ Activation function
- Backpropagation via traces
 - Efficient quantum backpropagation idea (2023)
- Optimization still classical
- Inherent noise more than in classical
- Epochs technically more due to measurements & shots

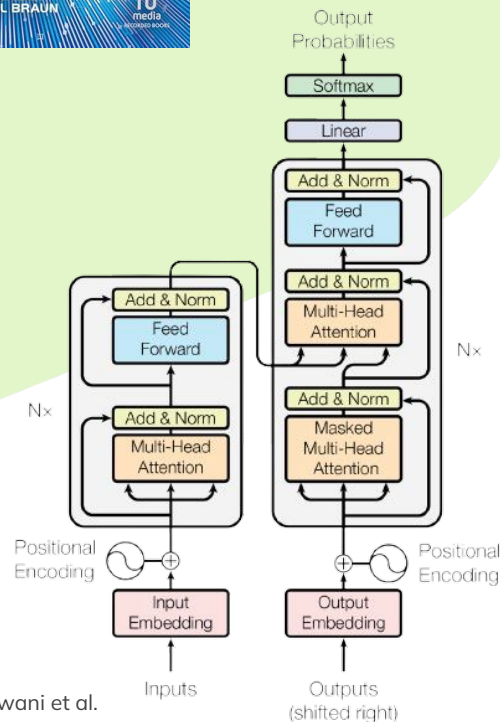


Evolution of neural networks

- [1958](#): Rosenblatt invents the Perceptron
 - Limited to linear separability
- [1969](#): Minsky & Papert love Perceptrons
 - Basically invent parallelism, but also cause 'AI Winter'
- 80s-90s: New techniques & models spur NNs again
 - Multilayer Perceptrons (MLPs), backpropagation, regularization
- [1990s](#): Nvidia GPUs perfect for neural networks
- 2010s: NN baby boom → DNNs, CNNs, RNNs, GANs...
- [2017](#): Google develops the Transformer
 - Self-attention mechanism allows dynamic focus of most relevant components.
- [2010s](#): Ethics and explainability of NNs (in AI) explored
- 2020s: ChatGPTs causes 'AI Summer'



Great
audiobook!



Credit: Vaswani et al.



doitnow
HPC Services

What's on the QNN menu?